

###Machine Learning Project

Laptop Price Prediction for SmartTech Co

1 Project Introduction

SmartTech Co. aims to strengthen its market position by understanding and predicting laptop prices using data-driven insights. With hundreds of models from various brands differing in RAM, CPU, storage, GPU, screen type, and operating systems, determining the right price is complex. Hence, this project develops a machine learning model that accurately predicts laptop prices and provides meaningful insights into the factors that influence pricing

2 Dataset Overview

The dataset (laptop.csv) contains detailed specifications of laptops collected from different brands. Each record represents a laptop configuration along with its price. | Column Name | Description | Type | | ----- | |

----- | | Company | Brand or manufacturer of the laptop | Categorical | | TypeName | Category of laptop (e.g., Ultrabook, Gaming, Notebook) | Categorical | | Inches | Screen size of the laptop in inches | Numerical | | ScreenResolution | Screen resolution and display type (e.g., 1920x1080 Touchscreen) | Text | | Cpu | Processor details (brand and model) | Text | | Ram | Size of the RAM in GB | Numerical | | Memory | Storage configuration (SSD, HDD, eMMC, Flash) | Text | | Gpu | Graphics processor (brand and model) | Text | | OpSys | Operating system of the laptop | Categorical | | Weight | Weight of the laptop in kilograms | Numerical | | Price | Target variable (price of the laptop) | Numerical |

3 Business Understanding

Business Problem

SmartTech Co. faces challenges in:

Accurately pricing laptops with diverse configurations.

Understanding the impact of specifications and brand on pricing.

Quickly evaluating new models' prices to stay competitive.

Business Objective

Build a predictive model to estimate laptop prices based on features.

Derive data insights to understand which features most influence pricing.

Enable real-time prediction for new models.

4 Machine Learning Project Workflow | Phase | Description | | ----- | |

----- | |
Phase 1: Data Understanding | Explore dataset shape, data types, missing values, and

distributions. || *Phase 2: Data Preprocessing* | Handle missing data, clean text, encode categorical variables, and normalize features. || *Phase 3: Feature Engineering* | Extract numeric and categorical insights like resolution, storage type, and processor brand. || *Phase 4: Model Development* | Train multiple models (Linear Regression, Random Forest, XGBoost). || *Phase 5: Model Evaluation* | Evaluate models using RMSE, MAE, and R^2 metrics. || *Phase 6: Hyperparameter Tuning* | Optimize the best-performing model using GridSearchCV. || *Phase 7: Insights & Interpretability* | Analyze feature importance and business takeaways. || *Phase 8: Deployment Ready Model* | Save the best model for real-time predictions. |

5 Success Criteria | Type | Description | Example Metric | | ----- |

----- | *Business Success* | Enable SmartTech Co. to predict prices accurately and gain competitive advantage. | Decision confidence & faster pricing | | *ML Success* | Achieve high R^2 and low RMSE on test data. | $R^2 \geq 0.75$, $RMSE \leq ₹20,000$ | | *Economic Success* | Reduce overpricing/underpricing losses. | Increased profit margins & market share |

6 Data Understanding & Insights

Initial data analysis showed:

Price increases with RAM, SSD, and screen resolution.

Gaming laptops and high PPI screens are priced higher.

Brand plays a significant role — Apple and MSI laptops tend to be premium.

Operating System: Windows-based laptops dominate the dataset.

Weight: Ultrabooks are lighter but more expensive due to design factors.

```
import pandas as pd
import numpy as np

df = pd.read_csv(r"/content/laptop.csv")
df.head()

{"summary": "{\n  \"name\": \"df\",\n  \"rows\": 1303,\n  \"fields\": [\n    {\n      \"column\": \"Unnamed: 0.1\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 376,\n        \"min\": 0,\n        \"max\": 1302,\n        \"num_unique_values\": 1303,\n        \"samples\": [\n          479,\n          1022,\n          298\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Unnamed: 0\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 376.4930271562809,\n        \"min\": 0.0,\n        \"max\": 1302.0,\n        \"num_unique_values\": 1273,\n        \"samples\": [\n          44.0,\n          1188.0,\n          133.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Company\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 19,\n        \"samples\": [\n          \"Apple\",\n          \"Lenovo\",\n          \"Xiaomi\"\n        ],\n
```

```

\"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"TypeName\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 6, \n      \"samples\": [ \n        \"Ultrabook\", \n        \"Notebook\", \n        \"Netbook\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"Inches\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 25, \n      \"samples\": [ \n        \"13\", \n        \"24\", \n        \"13.3\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"ScreenResolution\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 40, \n      \"samples\": [ \n        \"Touchscreen 1366x768\", \n        \"1600x900\", \n        \"4K Ultra HD / Touchscreen 3840x2160\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"Cpu\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 118, \n      \"samples\": [ \n        \"AMD A12-Series 9720P 2.7GHz\", \n        \"Intel Atom Z8350 1.92GHz\", \n        \"Intel Core i5 3.1GHz\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"Ram\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 10, \n      \"samples\": [ \n        \"24GB\", \n        \"16GB\", \n        \"64GB\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"Memory\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 40, \n      \"samples\": [ \n        \"256GB SSD + 500GB HDD\", \n        \"1.0TB Hybrid\", \n        \"64GB SSD\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"Gpu\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 110, \n      \"samples\": [ \n        \"AMD Radeon R7\", \n        \"AMD Radeon Pro 555\", \n        \"Intel Iris Plus Graphics 650\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"OpSys\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 9, \n      \"samples\": [ \n        \"Windows 7\", \n        \"No OS\", \n        \"Windows 10 S\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"Weight\", \n    \"properties\": { \n      \"dtype\": \"category\", \n      \"num_unique_values\": 189, \n      \"samples\": [ \n        \"2.21kg\", \n        \"3.74kg\", \n        \"1.65kg\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"Price\", \n    \"properties\": { \n      \"dtype\": \"number\", \n      \"std\": 37332.251004569865, \n      \"min\": 9270.72, \n      \"max\": 324954.72, \n      \"num_unique_values\": 777, \n

```

```

\"samples\": [\n          68464.8,\n          71341.92,\n          77788.8\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  }\n  \"type\": \"dataframe\", \"variable_name\": \"df\"}

# drop the unwanted columns
df.drop(columns=['Unnamed: 0', 'Unnamed: 0.1'], axis=1, inplace=True)
df.head()

{
  \"summary\": {
    \"name\": \"df\",
    \"rows\": 1303,
    \"fields\": [
      {
        \"column\": \"Company\",
        \"properties\": {
          \"dtype\": \"category\",
          \"num_unique_values\": 19,
          \"samples\": [
            \"Apple\",
            \"Lenovo\",
            \"Xiaomi\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        },
        \"column\": \"TypeName\",
        \"properties\": {
          \"dtype\": \"category\",
          \"num_unique_values\": 6,
          \"samples\": [
            \"Ultrabook\",
            \"Notebook\",
            \"Netbook\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        },
        \"column\": \"Inches\",
        \"properties\": {
          \"dtype\": \"category\",
          \"num_unique_values\": 25,
          \"samples\": [
            \"13\",
            \"24\",
            \"13.3\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        },
        \"column\": \"ScreenResolution\",
        \"properties\": {
          \"dtype\": \"category\",
          \"num_unique_values\": 40,
          \"samples\": [
            \"Touchscreen 1366x768\",
            \"1600x900\",
            \"4K Ultra HD / Touchscreen 3840x2160\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        },
        \"column\": \"Cpu\",
        \"properties\": {
          \"dtype\": \"category\",
          \"num_unique_values\": 118,
          \"samples\": [
            \"AMD A12-Series 9720P 2.7GHz\",
            \"Intel Atom Z8350 1.92GHz\",
            \"Intel Core i5 3.1GHz\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        },
        \"column\": \"Ram\",
        \"properties\": {
          \"dtype\": \"category\",
          \"num_unique_values\": 10,
          \"samples\": [
            \"24GB\",
            \"16GB\",
            \"64GB\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        },
        \"column\": \"Memory\",
        \"properties\": {
          \"dtype\": \"category\",
          \"num_unique_values\": 40,
          \"samples\": [
            \"256GB SSD + 500GB HDD\",
            \"1.0TB Hybrid\",
            \"64GB SSD\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        },
        \"column\": \"Gpu\",
        \"properties\": {
          \"dtype\": \"category\",
          \"num_unique_values\": 110,
          \"samples\": [
            \"AMD Radeon R7\",
            \"AMD Radeon Pro 555\",
            \"Intel Iris Plus Graphics 650\"
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        }
      ]
    }
  }
}

```

```

n    },\n    {\n        \"column\": \"OpSys\", \n        \"properties\": {\n            \"dtype\": \"category\", \n            \"num_unique_values\": 9, \n            \"samples\": [\n                \"Windows 7\", \n                \"No OS\", \n                \"Windows 10 S\" \n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        }, \n        {\n            \"column\": \"Weight\", \n            \"properties\": {\n                \"dtype\": \"category\", \n                \"num_unique_values\": 189, \n                \"samples\": [\n                    \"2.21kg\", \n                    \"3.74kg\", \n                    \"1.65kg\" \n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\" \n            }, \n            {\n                \"column\": \"Price\", \n                \"properties\": {\n                    \"dtype\": \"number\", \n                    \"std\": 37332.251004569865, \n                    \"min\": 9270.72, \n                    \"max\": 324954.72, \n                    \"num_unique_values\": 777, \n                    \"samples\": [\n                        68464.8, \n                        71341.92, \n                        77788.8 \n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\" \n                } \n            ] \n        } \n    ], \n    \"type\": \"dataframe\", \"variable_name\": \"df\"}

```

```

import numpy as np
# Replace "?" with np.nan
df = df.replace('?', np.nan)

```

```


import matplotlib.pyplot as plt
import seaborn as sns

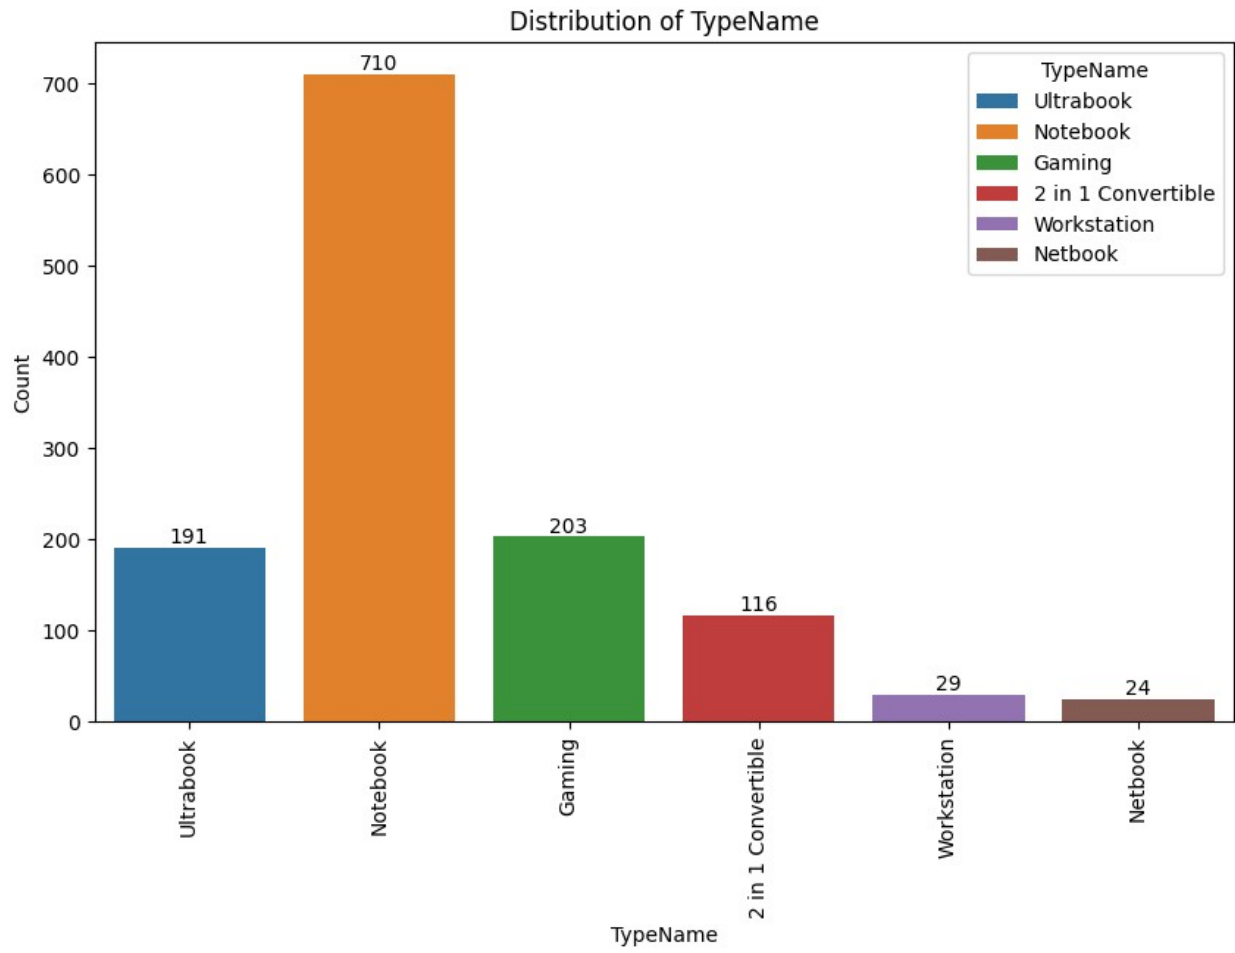
```

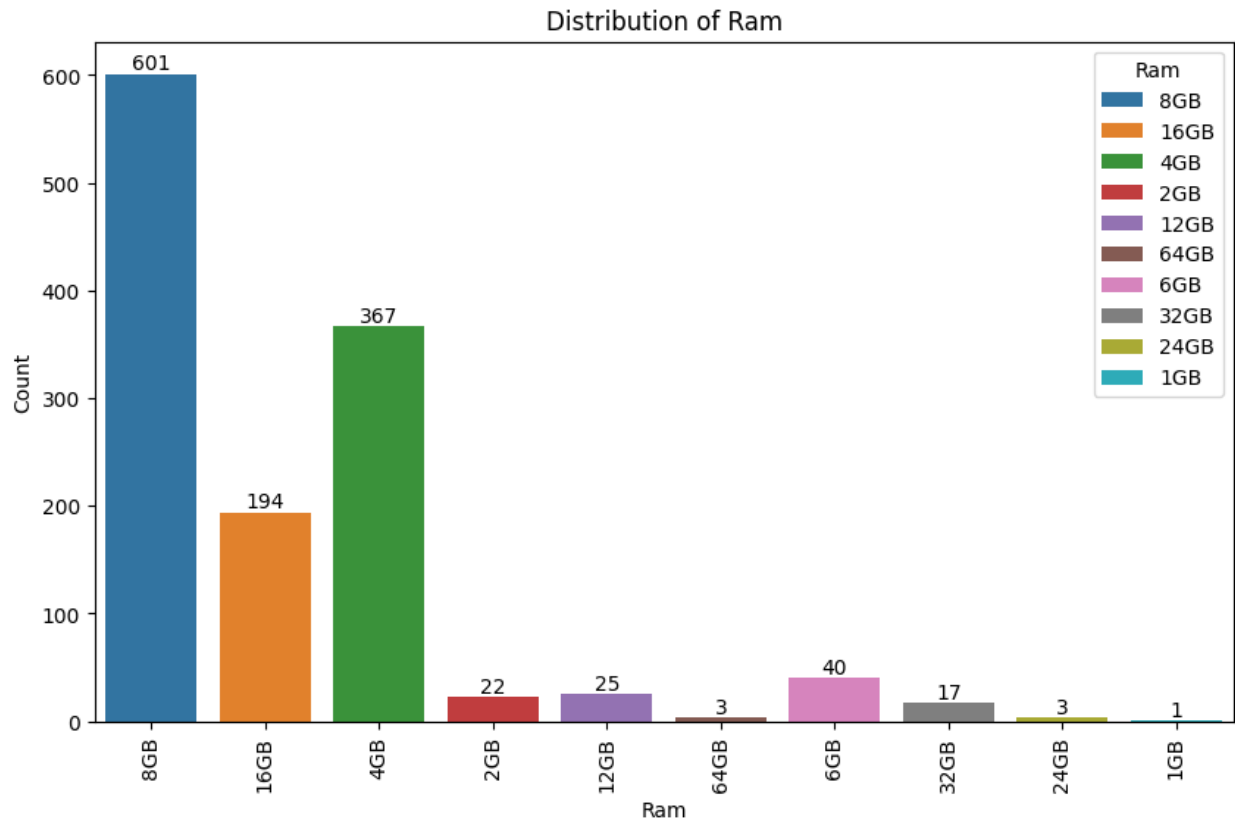
```

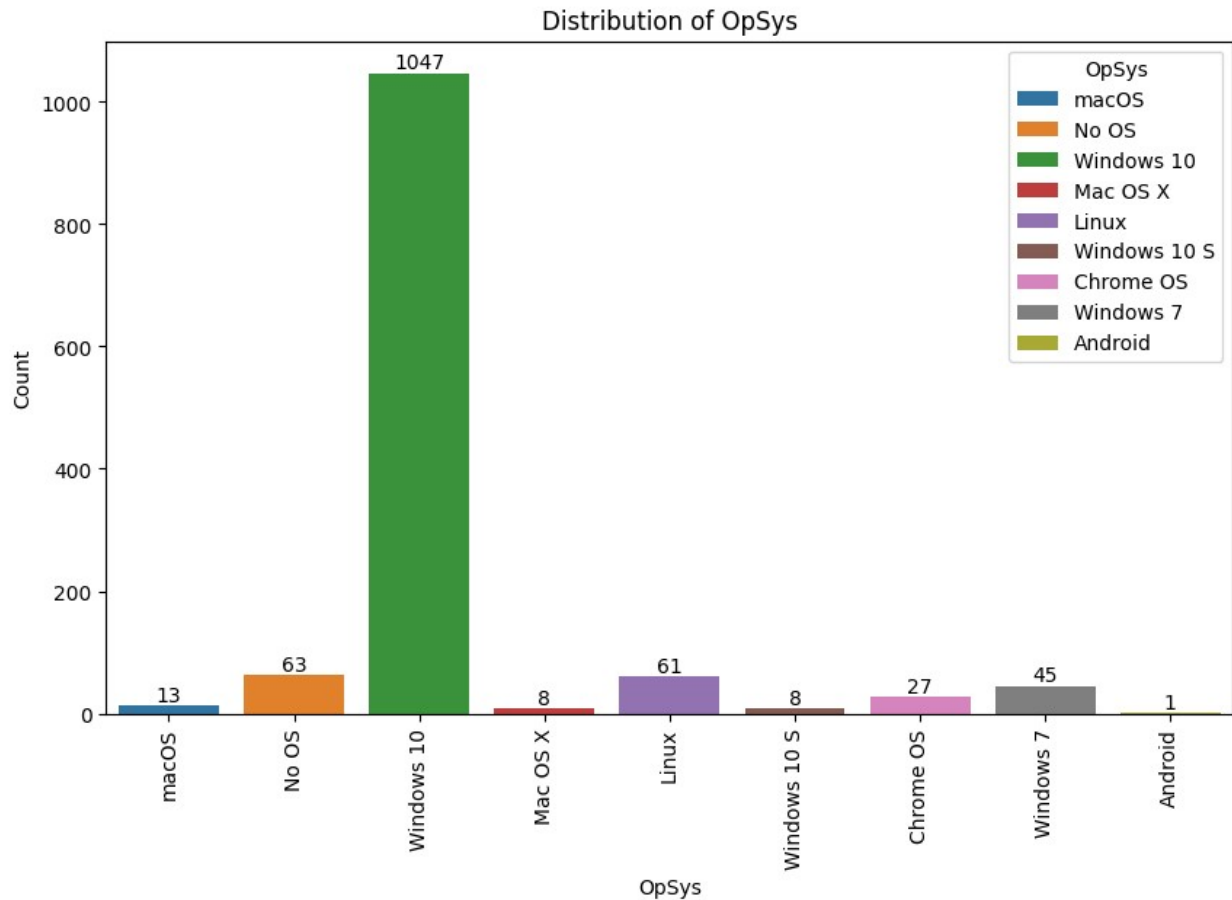
# Distribution of categorical variables
categorical_cols = ['Company', 'TypeName', 'Ram', 'OpSys']
for col in categorical_cols:
    plt.figure(figsize=(10, 6))
    ax = sns.countplot(x=col, data=df, hue=col)
    plt.title(f"Distribution of {col}")
    plt.xlabel(col)
    plt.ylabel("Count")
    plt.xticks(rotation=90)
    ## add data label
    for container in ax.containers:
        ax.bar_label(container)
    plt.show()

```









```
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
# Calculate the average price of laptops within each category
average_price_by_company = df.groupby('Company')
['Price'].mean().reset_index()
average_price_by_ram = df.groupby('Ram')['Price'].mean().reset_index()
average_price_by_inches = df.groupby('Inches')
['Price'].mean().reset_index()
average_price_by_inches = df.groupby('TypeName')
['Price'].mean().reset_index()
# Plot bar plots
plt.figure(figsize=(12, 8))
sns.barplot(x='Company', y='Price', data=df, estimator=lambda x:
sum(x) / len(x), palette='viridis')
plt.title('Average Laptop Prices by Company')
plt.xticks(rotation=45)
plt.show()

plt.figure(figsize=(12, 8))
sns.barplot(x='Ram', y='Price', data=df, estimator=lambda x: sum(x) /
```

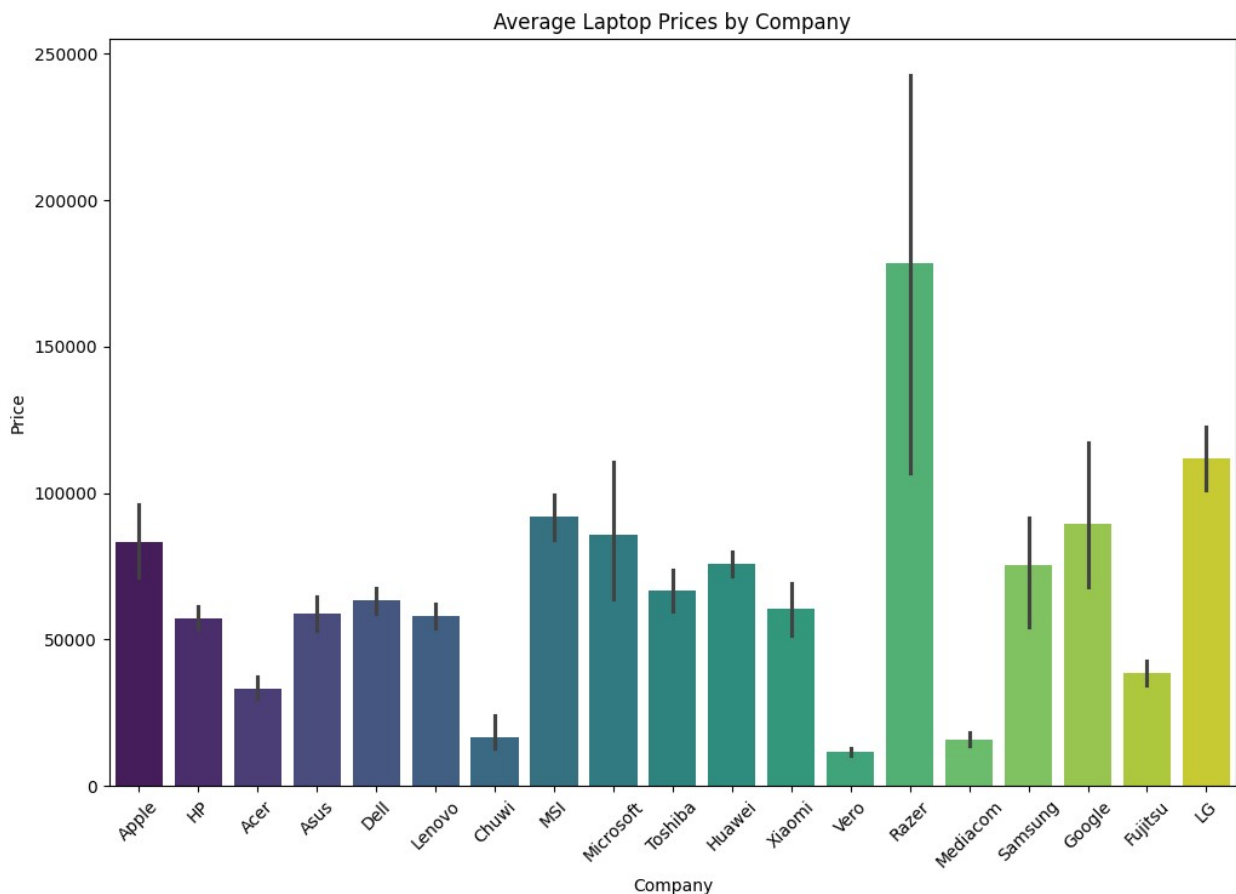
```

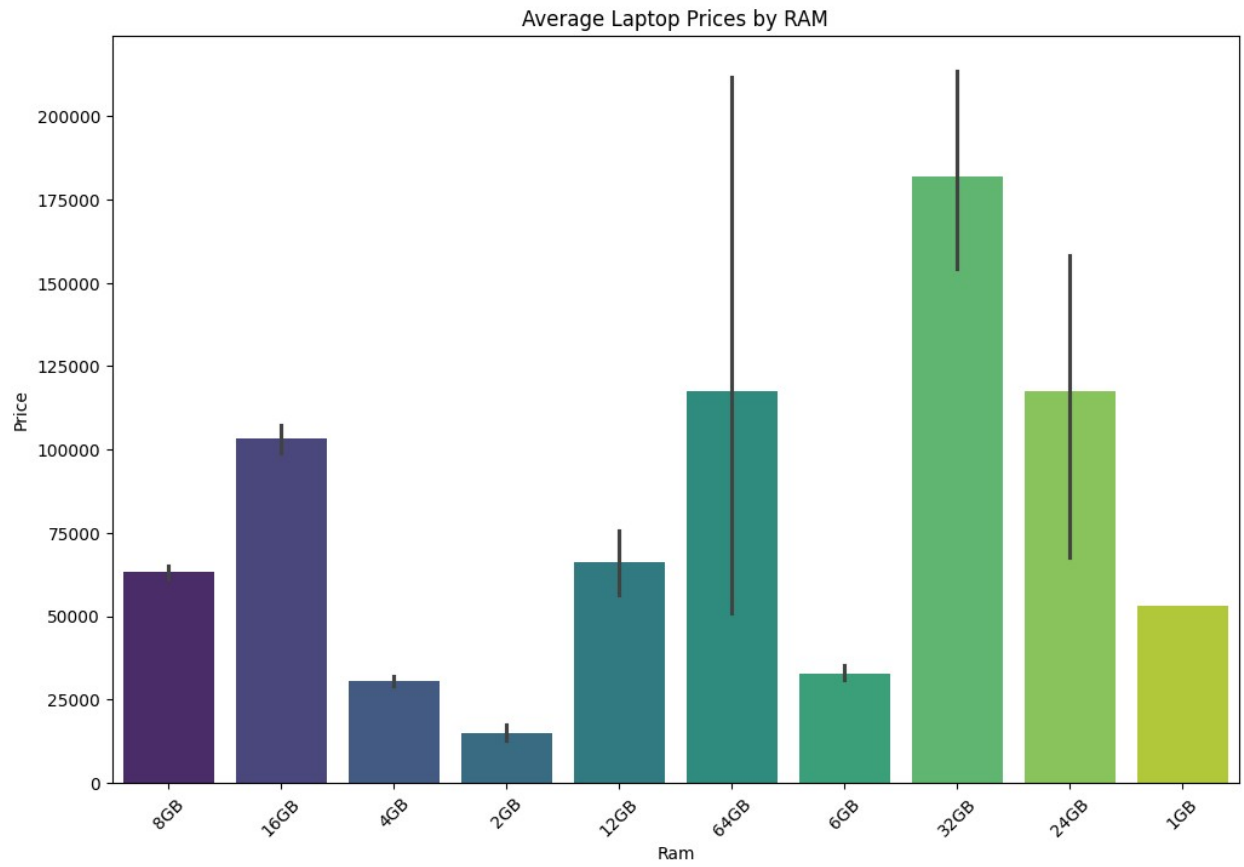
len(x),palette='viridis')
plt.title('Average Laptop Prices by RAM')
plt.xticks(rotation=45)
plt.show()

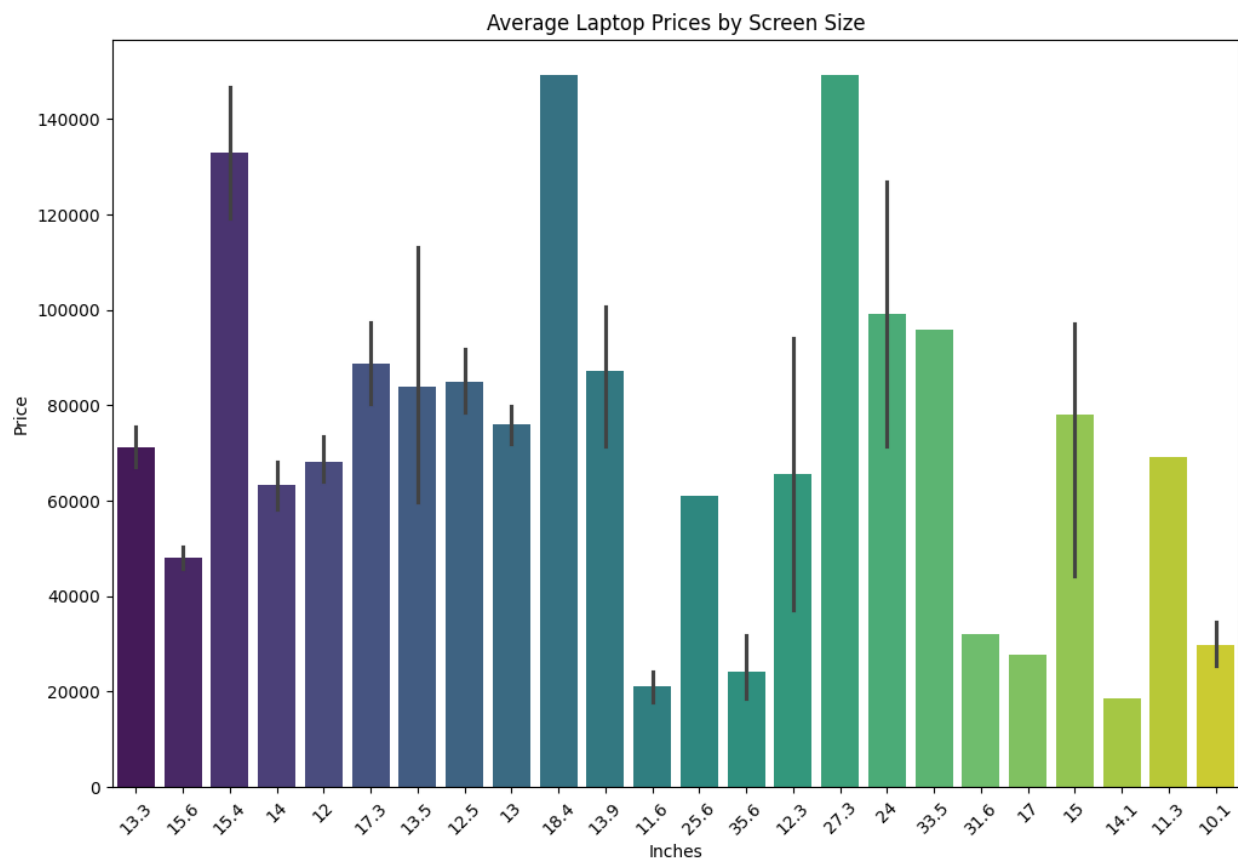
plt.figure(figsize=(12, 8))
sns.barplot(x='Inches', y='Price', data=df, estimator=lambda x: sum(x) / len(x),palette='viridis')
plt.title('Average Laptop Prices by Screen Size')
plt.xticks(rotation=45)
plt.show()

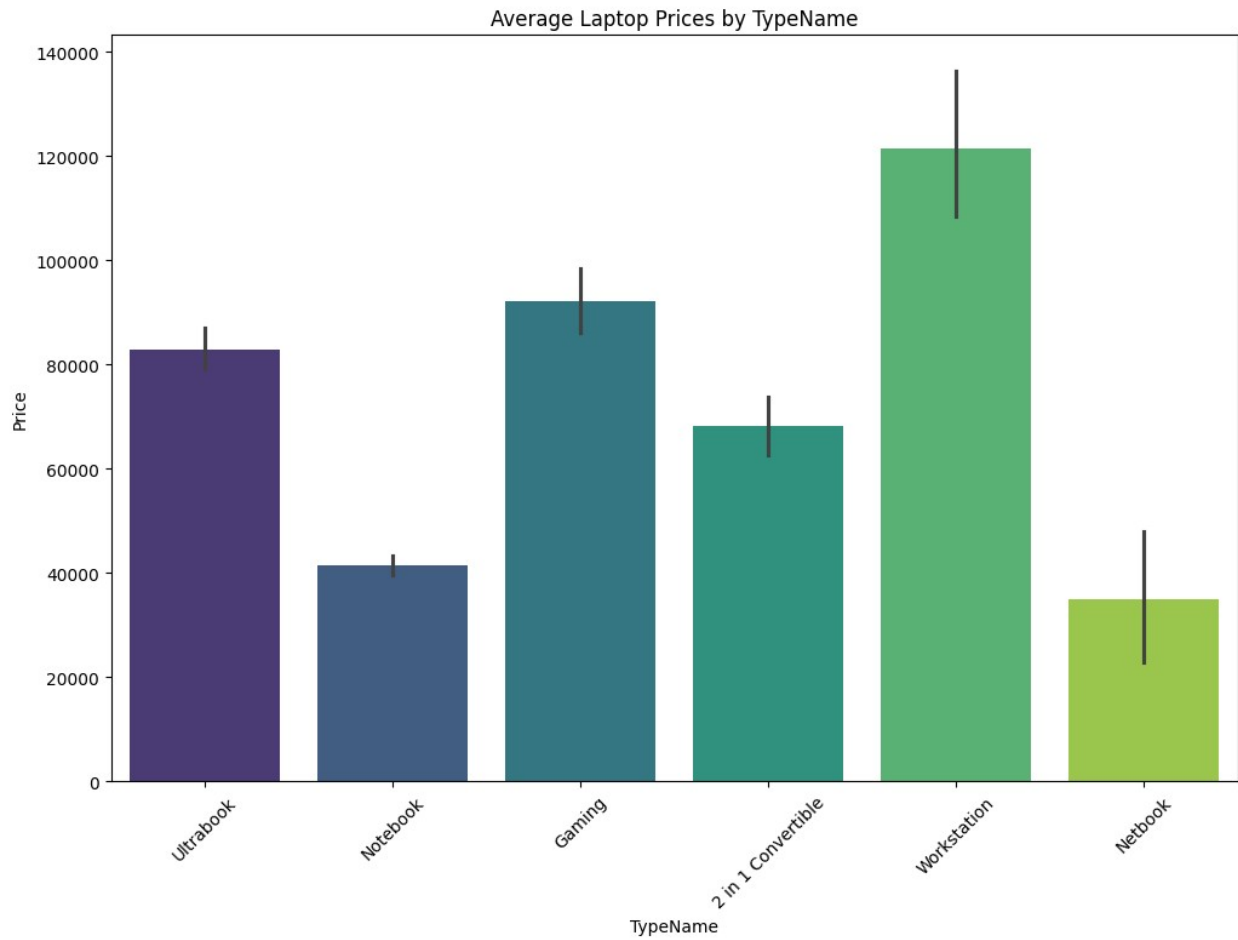
plt.figure(figsize=(12, 8))
sns.barplot(x='TypeName', y='Price', data=df, estimator=lambda x: sum(x) / len(x),palette='viridis')
plt.title('Average Laptop Prices by TypeName')
plt.xticks(rotation=45)
plt.show()

```









###Checking null values

```
df.isnull().sum()
```

```
Company      30
TypeName     30
Inches       31
ScreenResolution  30
Cpu          30
Ram          30
Memory       31
Gpu          30
OpSys        30
Weight       31
Price        30
dtype: int64
```

###Dropped null values and checking after imputation

```
df.dropna(inplace=True)
print(df.isnull().sum())
```

```
Company          0
TypeName         0
Inches           0
ScreenResolution 0
Cpu              0
Ram              0
Memory           0
Gpu              0
OpSys            0
Weight           0
Price            0
dtype: int64
```

###Feature Engineering

```
import pandas as pd
import numpy as np
# Define the function to extract CPU series
def extract_series(cpu):
    if pd.isna(cpu):
        return 'Unknown'
    if 'Intel Core i3' in cpu:
        return 'Intel Core i3'
    elif 'Intel Core i5' in cpu:
        return 'Intel Core i5'
    elif 'Intel Core i7' in cpu:
        return 'Intel Core i7'
    elif 'Intel Core M' in cpu:
        return 'Intel Core M'
    elif 'Intel Atom' in cpu:
        return 'Intel Atom'
    elif 'Intel Pentium' in cpu:
        return 'Intel Pentium'
    elif 'Intel Celeron' in cpu:
        return 'Intel Celeron'
    elif 'Intel Xeon' in cpu:
        return 'Intel Xeon'
    elif 'AMD E-Series' in cpu:
        return 'AMD E-Series'
    elif 'AMD A6-Series' in cpu:
        return 'AMD A6-Series'
    elif 'AMD A9-Series' in cpu:
        return 'AMD A9-Series'
    elif 'AMD A10-Series' in cpu:
```

```

        return 'AMD A10-Series'
    elif 'AMD A12-Series' in cpu:
        return 'AMD A12-Series'
    elif 'AMD FX' in cpu:
        return 'AMD FX'
    elif 'AMD Ryzen' in cpu:
        return 'AMD Ryzen'
    else:
        return 'Other'

# Apply the function to the 'Cpu' column and create a new column
'Cpu_Series'
df['Cpu_Series'] = df['Cpu'].apply(extract_series)

df['Touchscreen'] = df['ScreenResolution'].apply(lambda x:
'Touchscreen' if 'Touchscreen' in x else 'NonTouchscreen')
df['Touchscreen'].value_counts()

Touchscreen
NonTouchscreen    1085
Touchscreen        185
Name: count, dtype: int64

import pandas as pd

def categorize_os(os_name):
    if 'Windows 10' in os_name or 'Windows 7' in os_name or 'Windows
10 S' in os_name:
        return 'Windows'
    elif 'No OS' in os_name:
        return 'No OS'
    elif 'Mac OS X' in os_name or 'macOS' in os_name:
        return 'macOS'
    elif 'Linux' in os_name or 'Chrome OS' in os_name or 'Android' in
os_name:
        return 'other/Linux'
    else:
        return 'Unknown'

# Assuming 'df' is your DataFrame and 'OS' is the column containing
the OS names
df['OS'] = df['OpSys'].apply(categorize_os)

# Verify the result
print(df[['OpSys', 'OS']].head())
print(df['OS'].value_counts())

```

	OpSys	OS
0	macOS	macOS
1	macOS	macOS
2	No OS	No OS

```

3  macOS  macOS
4  macOS  macOS
OS
Windows      1097
other/Linux   89
No OS         63
macOS         21
Name: count, dtype: int64

```

```

def categorize_gpu(gpu_name):
    if 'Intel' in gpu_name:
        return 'Intel'
    elif 'AMD' in gpu_name:
        return 'AMD'
    elif 'Nvidia' in gpu_name or 'NVIDIA' in gpu_name:
        return 'Nvidia'
    elif 'ARM Mali' in gpu_name:
        return 'ARM Mali'
    else:
        return 'other'

```

Apply the function to the 'Gpu' column and create a new column 'GpuCategory'

```
df['GpuCategory'] = df['Gpu'].apply(categorize_gpu)
```

Verify the result

```

print(df[['Gpu', 'GpuCategory']].head())
print(df['GpuCategory'].value_counts())

```

	Gpu	GpuCategory
0	Intel Iris Plus Graphics 640	Intel
1	Intel HD Graphics 6000	Intel
2	Intel HD Graphics 620	Intel
3	AMD Radeon Pro 455	AMD
4	Intel Iris Plus Graphics 650	Intel

```

GpuCategory
Intel      702
Nvidia     392
AMD        175
ARM Mali     1
Name: count, dtype: int64

```

```
import re
```

Function to extract SSD and HDD storage from memory string

```

def extract_storage(memory):
    ssd = 0
    hdd = 0
    flash_storage = 0
    hybrid = 0

```



```

    ssd_match = re.findall(r'(\d+(?:\.\d+)?)(?=GB SSD|TB SSD)',
memory)
    hdd_match = re.findall(r'(\d+(?:\.\d+)?)(?=GB HDD|TB HDD)',
memory)
    flash_storage_match = re.findall(r'(\d+(?:\.\d+)?)(?=GB Flash
Storage|TB Flash Storage)', memory)
    hybrid_match = re.findall(r'(\d+(?:\.\d+)?)(?=GB Hybrid|TB
Hybrid)', memory)

    if ssd_match:
        ssd = sum(int(float(x.replace('TB', '1024').replace('GB',
''))) for x in ssd_match)

    if hdd_match:
        hdd = sum(int(float(x.replace('TB', '1024').replace('GB',
''))) for x in hdd_match)

    if flash_storage_match:
        flash_storage = sum(int(float(x.replace('TB',
'1024').replace('GB', ''))) for x in flash_storage_match)

    if hybrid_match:
        hybrid = sum(int(float(x.replace('TB', '1024').replace('GB',
''))) for x in hybrid_match)

    return pd.Series([ssd, hdd, flash_storage, hybrid])

# Apply the function to the DataFrame and create new columns
df[['SSD', 'HDD', 'Flash Storage', 'Hybrid']] =
df['Memory'].apply(extract_storage)

# Print the result
print(df[['Memory', 'SSD', 'HDD', 'Flash
Storage', 'Hybrid']].value_counts())

```

Memory	SSD	HDD	Flash Storage	Hybrid	
256GB SSD	256	0	0	0	399
1TB HDD	0	1	0	0	217
500GB HDD	0	500	0	0	130
512GB SSD	512	0	0	0	116
128GB SSD + 1TB HDD	128	1	0	0	92
128GB SSD	128	0	0	0	74
256GB SSD + 1TB HDD	256	1	0	0	71
32GB Flash Storage	0	0	32	0	37
2TB HDD	0	2	0	0	16
64GB Flash Storage	0	0	64	0	14
512GB SSD + 1TB HDD	512	1	0	0	14
1TB SSD	1	0	0	0	13
256GB SSD + 2TB HDD	256	2	0	0	10
1.0TB Hybrid	0	0	0	1	9

256GB Flash Storage	0	0	256	0	8
16GB Flash Storage	0	0	16	0	7
32GB SSD	32	0	0	0	6
180GB SSD	180	0	0	0	4
128GB Flash Storage	0	0	128	0	4
512GB SSD + 2TB HDD	512	2	0	0	3
16GB SSD	16	0	0	0	3
128GB SSD + 2TB HDD	128	2	0	0	2
1TB SSD + 1TB HDD	1	1	0	0	2
512GB Flash Storage	0	0	512	0	2
256GB SSD + 500GB HDD	256	500	0	0	2
256GB SSD + 256GB SSD	512	0	0	0	2
1.0TB HDD	0	1	0	0	1
128GB HDD	0	128	0	0	1
256GB SSD + 1.0TB Hybrid	256	0	0	1	1
1TB HDD + 1TB HDD	0	2	0	0	1
240GB SSD	240	0	0	0	1
512GB SSD + 1.0TB Hybrid	512	0	0	1	1
508GB Hybrid	0	0	0	508	1
32GB HDD	0	32	0	0	1
512GB SSD + 256GB SSD	768	0	0	0	1
512GB SSD + 512GB SSD	1024	0	0	0	1
64GB Flash Storage + 1TB HDD	0	1	64	0	1
64GB SSD	64	0	0	0	1
8GB SSD	8	0	0	0	1

Name: count, dtype: int64

```

import matplotlib.pyplot as plt
import seaborn as sns

avg_price_cpu_series = df.groupby('Cpu_Series')
['Price'].mean().reset_index()
avg_price_gpu_category = df.groupby('GpuCategory')
['Price'].mean().reset_index()
avg_price_OS = df.groupby('OS')['Price'].mean().reset_index()
avg_price_touchscreen = df.groupby('Touchscreen')
['Price'].mean().reset_index()

# Melt the DataFrame to have 'StorageType' and 'Price' columns
melted_df = df.melt(id_vars='Price', value_vars=['SSD', 'HDD', 'Flash Storage', 'Hybrid'], var_name='StorageType', value_name='Count')
melted_df = melted_df[melted_df['Count'] > 0] # Filter out rows with count 0

# Group by 'StorageType' and calculate the mean price
avg_prices_storage_types = melted_df.groupby('StorageType')
['Price'].mean().reset_index()

# Plot the average price for each storage type
plt.figure(figsize=(10, 6))

```

```

sns.barplot(data=avg_prices_storage_types, x='StorageType',
y='Price',palette='viridis')
plt.title('Average Price by Storage Type')
plt.xlabel('Storage Type')
plt.ylabel('Average Price')
plt.xticks(rotation=45)
plt.show()

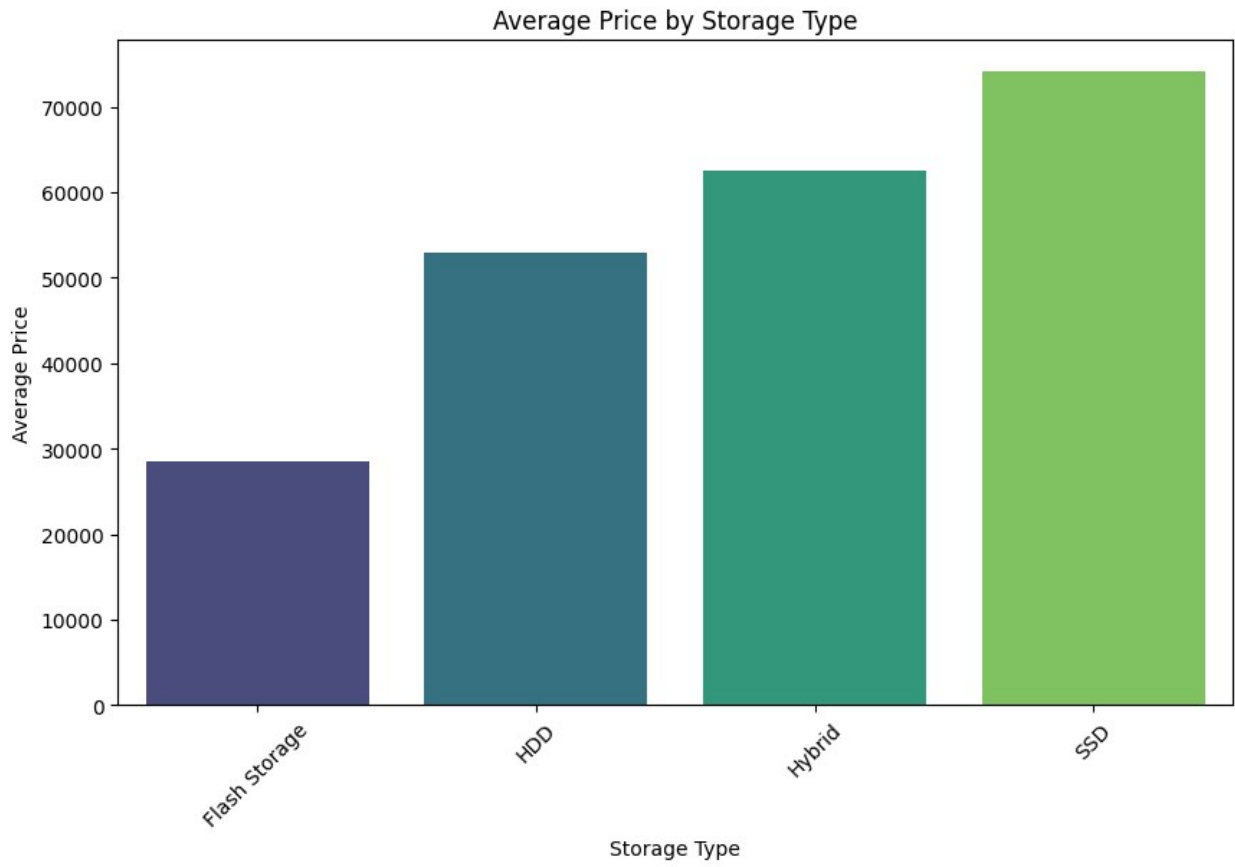
# Plot average price by CPU Series
plt.figure(figsize=(8, 6))
sns.barplot(data=avg_price_cpu_series, x='Cpu_Series',
y='Price',palette='viridis')
plt.title('Average Price by CPU Series')
plt.xlabel('CPU Series')
plt.ylabel('Average Price')
plt.xticks(rotation=90)
plt.show()


# Plot average price by GPU Category
plt.figure(figsize=(8, 6))
sns.barplot(data=avg_price_gpu_category, x='GpuCategory',
y='Price',palette='viridis')
plt.title('Average Price by GPU Category')
plt.xlabel('GPU Category')
plt.ylabel('Average Price')
plt.show()


# Plotting
plt.figure(figsize=(10, 6))
sns.barplot(x='OS', y='Price', data=avg_price_OS,palette='viridis')
plt.title('Average Price of Devices by Operating System')
plt.xlabel('Operating System')
plt.ylabel('Average Price')
plt.xticks(rotation=45)
plt.show()

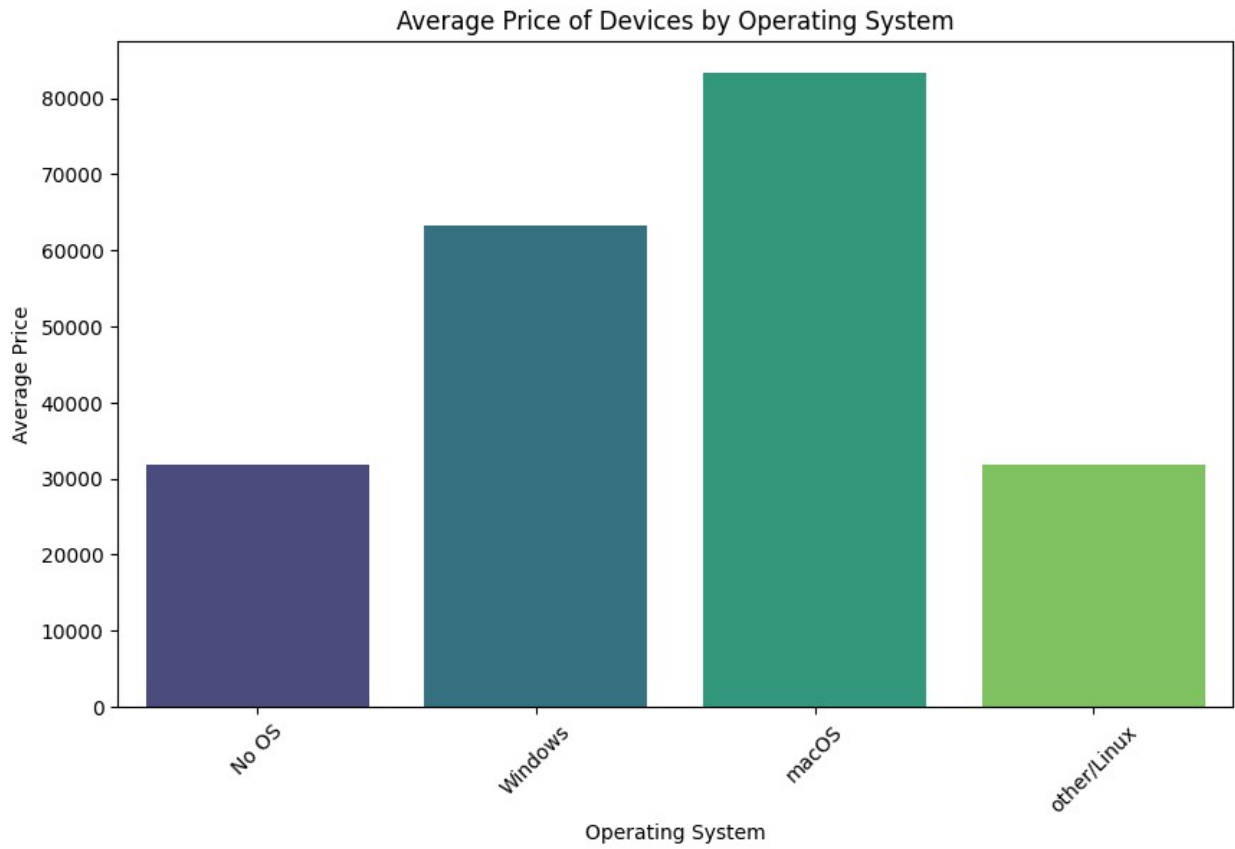
# Plot average price by Touchscreen
plt.figure(figsize=(8, 6))
sns.barplot(data=avg_price_touchscreen, x='Touchscreen',
y='Price',palette='viridis')
plt.title('Average Price by Touchscreen')
plt.xlabel('Touchscreen')
plt.ylabel('Average Price')
plt.show()

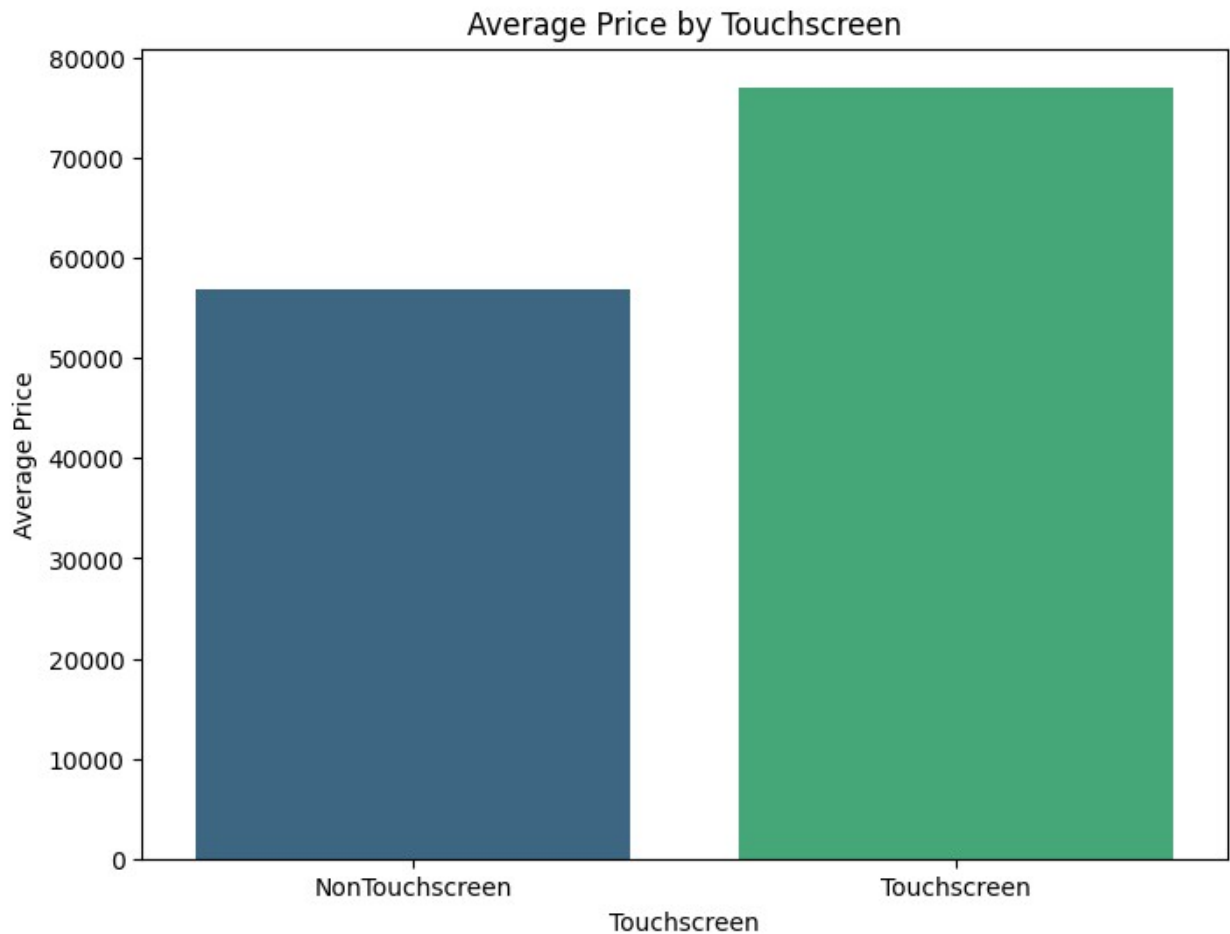
```











```
# Function to extract X and Y resolutions and panel type using regular
expressions
def extract_resolution(resolution):
    match = re.match(r'(?:(?P<panel_type>.*?)(?:\sPanel)?\s?(?:\s/)?\s?
(?:Touchscreen)?\s?)(?P<x_res>\d+)x(?:P<y_res>\d+)', resolution)
    if match:
        panel_type = match.group('panel_type') or 'Unknown'
        if panel_type == "IPS":
            panel_type = "IPS Panel"
        x_res = int(match.group('x_res'))
        y_res = int(match.group('y_res'))
        return pd.Series([panel_type, x_res, y_res])
    else:
        return pd.Series(['Unknown', 0, 0])

# Apply the function to the 'ScreenResolution' column and create new
columns 'PanelType', 'X_res', and 'Y_res'
df[['PanelType', 'X_res', 'Y_res']] =
df['ScreenResolution'].apply(extract_resolution)

# Function to modify the PanelType column
```



```
def modify_panel_type(panel_type):
    panel_type = re.sub(r'(?i)(?:\\|\\|Touchscreen)', '', panel_type)
# Remove "/", "\", and "Touchscreen"
    if panel_type.startswith("IPS Panel"):
        panel_type = "IPS Panel"
    panel_type = re.sub(r'^\s+|\s+$', '', panel_type) # Strip leading
and trailing whitespaces
    return panel_type

# Apply the modification to the PanelType column
df['PanelType'] = df['PanelType'].apply(modify_panel_type)

# Print the result
print(df[['ScreenResolution', 'PanelType', 'X_res', 'Y_res']])
```

		ScreenResolution	PanelType	X_res
Y_res				
0	IPS Panel Retina Display	2560x1600	IPS Panel	2560
1600				
1		1440x900	Unknown	1440
900				
2	Full HD	1920x1080	Full HD	1920
1080				
3	IPS Panel Retina Display	2880x1800	IPS Panel	2880
1800				
4	IPS Panel Retina Display	2560x1600	IPS Panel	2560
1600				
...	
...				
1298	IPS Panel Full HD / Touchscreen	1920x1080	IPS Panel	1920
1080				
1299	IPS Panel Quad HD+ / Touchscreen	3200x1800	IPS Panel	3200
1800				
1300		1366x768	Unknown	1366
768				
1301		1366x768	Unknown	1366
768				
1302		1366x768	Unknown	1366
768				

[1270 rows x 4 columns]

###Label Encoding

```
from sklearn.preprocessing import LabelEncoder

columns_to_encode =
['Company', 'TypeName', 'OS', 'Touchscreen', 'Cpu_Series', 'GpuCategory', 'P
anelType']
```

```

# Dictionary to store label encoders
label_encoders = {}

# Label encode each column in the list
for col in columns_to_encode:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

# To print the mapping for each column
for col, le in label_encoders.items():
    print(f"Mapping for {col}:")
    for class_index, class_label in enumerate(le.classes_):
        print(f"    {class_label}: {class_index}")

```

Mapping for Company:

```

Acer: 0
Apple: 1
Asus: 2
Chuwi: 3
Dell: 4
Fujitsu: 5
Google: 6
HP: 7
Huawei: 8
LG: 9
Lenovo: 10
MSI: 11
Mediacom: 12
Microsoft: 13
Razer: 14
Samsung: 15
Toshiba: 16
Vero: 17
Xiaomi: 18

```

Mapping for TypeName:

```

2 in 1 Convertible: 0
Gaming: 1
Netbook: 2
Notebook: 3
Ultrabook: 4
Workstation: 5

```

Mapping for OS:

```

No OS: 0
Windows: 1
macOS: 2
other/Linux: 3

```

Mapping for Touchscreen:

```

NonTouchscreen: 0

```

```

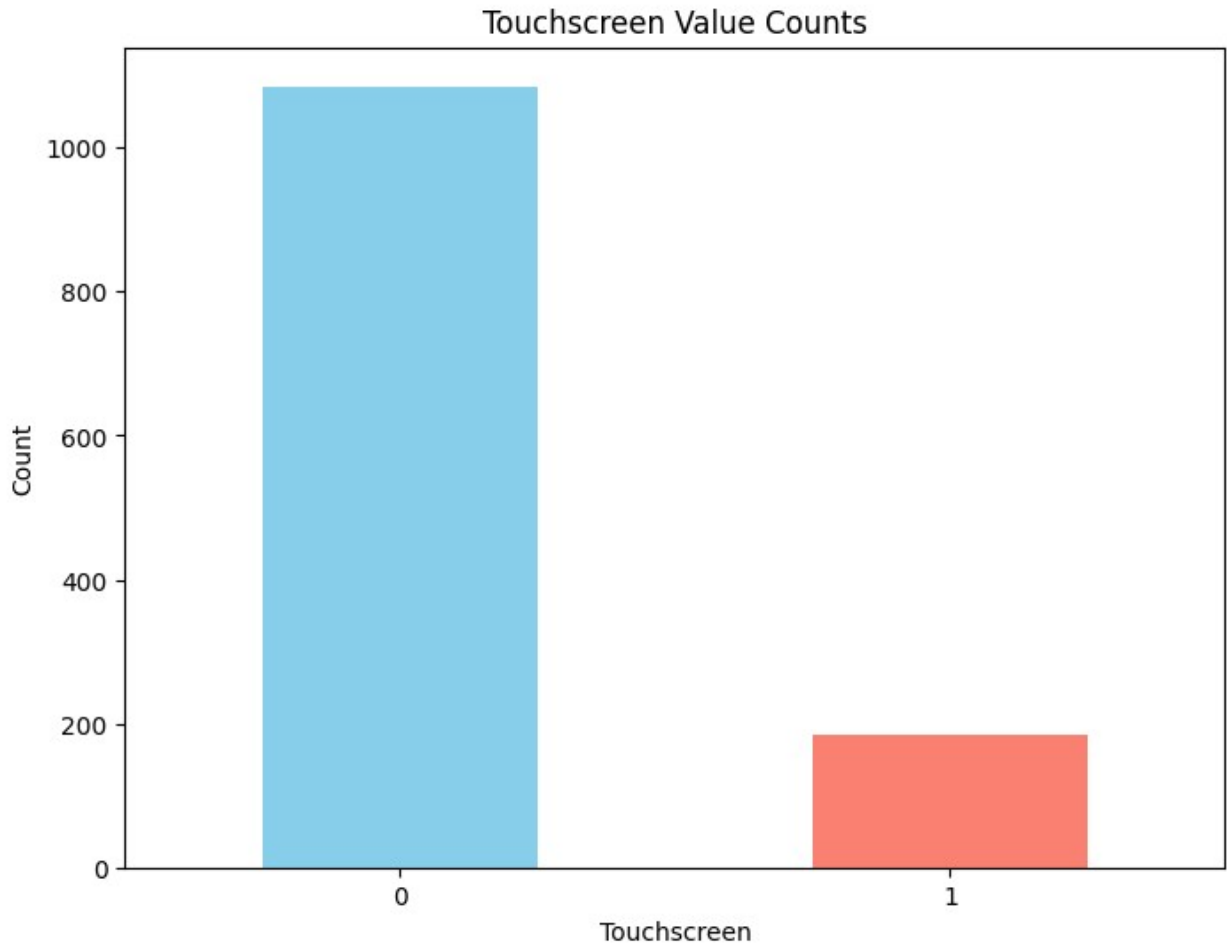
    Touchscreen: 1
Mapping for Cpu_Series:
    AMD A10-Series: 0
    AMD A12-Series: 1
    AMD A6-Series: 2
    AMD A9-Series: 3
    AMD E-Series: 4
    AMD FX: 5
    AMD Ryzen: 6
    Intel Atom: 7
    Intel Celeron: 8
    Intel Core M: 9
    Intel Core i3: 10
    Intel Core i5: 11
    Intel Core i7: 12
    Intel Pentium: 13
    Intel Xeon: 14
    Other: 15
Mapping for GpuCategory:
    AMD: 0
    ARM Mali: 1
    Intel: 2
    Nvidia: 3
Mapping for PanelType:
    4K Ultra HD: 0
    Full HD: 1
    IPS Panel: 2
    Quad HD+: 3
    Unknown: 4

import pandas as pd
import matplotlib.pyplot as plt

# Get the value counts
value_counts = df['Touchscreen'].value_counts()

# Plotting
plt.figure(figsize=(8, 6))
value_counts.plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Touchscreen Value Counts')
plt.xlabel('Touchscreen')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.show()

```



```
# Convert 'Weight' column to string type
df['Weight'] = df['Weight'].astype(str)
df['Ram'] = df['Ram'].astype(str)

# Extract only numeric values from 'Weight' column
df['Weight'] = df['Weight'].str.extract(r'([\d.]+)').astype(float)
df['Ram'] = df['Ram'].str.extract(r'(\d+)').astype(float)
df

{"type": "dataframe", "variable_name": "df"}

df.columns
Index(['Company', 'TypeName', 'Inches', 'ScreenResolution', 'Cpu',
      'Ram',
      'Memory', 'Gpu', 'OpSys', 'Weight', 'Price', 'Cpu_Series',
      'Touchscreen', 'OS', 'GpuCategory', 'SSD', 'HDD', 'Flash
Storage',
      'Hybrid', 'PanelType', 'X_res', 'Y_res'],
      dtype='object')
```

###Selecting features and target variable

```
x=df[['Company','TypeName','Ram','Inches','OS','Cpu_Series','Touchscreen','GpuCategory','SSD','HDD','FlashStorage','Hybrid','PanelType','X_res','Y_res']]
```

x

```
{"summary":{"\n  \"name\": \"x\",\n  \"rows\": 1270,\n  \"fields\": [\n    {\n      \"column\": \"Company\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 4,\n        \"min\": 0,\n        \"max\": 18,\n        \"num_unique_values\": 19,\n        \"samples\": [\n          1,\n          10,\n          18\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"TypeName\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1,\n        \"min\": 0,\n        \"max\": 5,\n        \"num_unique_values\": 6,\n        \"samples\": [\n          4,\n          3,\n          2\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Ram\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 5.566929534499876,\n        \"min\": 1.0,\n        \"max\": 64.0,\n        \"num_unique_values\": 10,\n        \"samples\": [\n          24.0,\n          16.0,\n          64.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Inches\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 24,\n        \"samples\": [\n          \"13\",\n          \"24\",\n          \"13.3\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"OS\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 3,\n        \"num_unique_values\": 4,\n        \"samples\": [\n          0,\n          3,\n          2\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Cpu_Series\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2,\n        \"min\": 0,\n        \"max\": 15,\n        \"num_unique_values\": 16,\n        \"samples\": [\n          11,\n          12,\n          4\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Touchscreen\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          1,\n          0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"GpuCategory\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 3,\n        \"num_unique_values\": 4,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}
```


###Splitting the data into training and testing sets

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y, test_size= 0.2,
random_state = 42)
print("train", x_train.shape)
print("test", x_test.shape)

train (1016, 15)
test (254, 15)

from sklearn.preprocessing import StandardScaler

# Feature scaling
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

x_train['Inches'] = pd.to_numeric(x_train['Inches'], errors='coerce')
x_test['Inches'] = pd.to_numeric(x_test['Inches'], errors='coerce')

from sklearn.linear_model import LinearRegression
import xgboost as xgb
from sklearn.ensemble import
RandomForestRegressor,GradientBoostingRegressor
from sklearn.metrics import r2_score,mean_squared_error

models = {
    'Linear Regression': LinearRegression(),
    'Random Forest Regression':
RandomForestRegressor(random_state=42),
    'Gradient Boosting Regressor' :
GradientBoostingRegressor(random_state=42),
    'XGBoost Regressor': xgb.XGBRegressor(random_state=42)
}

results = {}
for name, model in models.items():
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    r2 = r2_score(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    results[name] = {'R_2 Score': r2, 'Mean Squared Error': mse ,
'Predictions': y_pred}

/usr/local/lib/python3.12/dist-packages/sklearn/base.py:1389:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples,), for
example using ravel().
```

```
    return fit_method(estimator, *args, **kwargs)
/usr/local/lib/python3.12/dist-packages/sklearn/ensemble/_gb.py:672:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
```

```
    y = column_or_1d(y, warn=True) # TODO: Is this still required?
```

```
# Printing results
```

```
for name, metrics in results.items():
    print(f"\nModel: {name}")
    print(f"R-squared: {metrics['R_2 Score']}")
    print(f"Mean Squared Error: {metrics['Mean Squared Error']}\n")
```

```
Model: Linear Regression
R-squared: 0.6039912817054862
Mean Squared Error: 440325039.99111944
```

```
Model: Random Forest Regression
R-squared: 0.8338655456499053
Mean Squared Error: 184726135.7544251
```

```
Model: Gradient Boosting Regressor
R-squared: 0.8126058891921002
Mean Squared Error: 208364905.9317487
```

```
Model: XGBoost Regressor
R-squared: 0.7834740281105042
Mean Squared Error: 240756832.0
```

Hyperparameter Tuning

```
import pandas as pd
from sklearn.linear_model import LinearRegression
import xgboost as xgb
import pickle
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.model_selection import RandomizedSearchCV

models = {
    'Linear Regression': LinearRegression(),
    'Random Forest Regression':
    RandomForestRegressor(random_state=42),
```



```

    'Gradient Boosting Regressor':
GradientBoostingRegressor(random_state=42),
    'XGBoost Regressor': xgb.XGBRegressor(random_state=42)
}

# Define parameter distributions for each model
param_dists = {
    'Linear Regression': {
        'fit_intercept': [True, False],
        'copy_X': [True, False],
        'positive': [True, False]
    },
    'Random Forest Regression': {
        'n_estimators': [100, 200, 300],
        'max_depth': [None, 10, 20, 30],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4],
        'max_features': ['auto', 'sqrt']
    },
    'Gradient Boosting Regressor': {
        'n_estimators': [100, 200, 300, 400, 500],
        'learning_rate': [0.01, 0.05, 0.1, 0.2],
        'max_depth': [3, 4, 5, 6],
        'subsample': [0.7, 0.8, 0.9, 1.0],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    },
    'XGBoost Regressor': {
        'n_estimators': [100, 200, 300],
        'learning_rate': [0.01, 0.1, 0.05],
        'max_depth': [3, 4, 5],
        'subsample': [0.8, 0.9, 1.0],
        'colsample_bytree': [0.8, 0.9, 1.0]
    }
}

# Perform RandomizedSearchCV for each model
best_models = {}
for name, model in models.items():
    print(f"Tuning {name}...")
    random_search = RandomizedSearchCV(estimator=model,
param_distributions=param_dists[name], cv=5, n_iter=10, scoring='r2',
n_jobs=-1, random_state=42)
    random_search.fit(x_train, y_train)
    best_models[name] = random_search.best_estimator_
    print(f"Best parameters for {name}: {random_search.best_params_}")
    print(f"Best R-squared for {name}: {random_search.best_score_}")

# Evaluate the best models on the test set
results = {}

```

```

best_model_name = None
best_r2_score = -float('inf')

for name, model in best_models.items():
    y_pred = model.predict(x_test)
    r2 = r2_score(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    results[name] = {'R_2 Score': r2, 'Mean Squared Error': mse,
                    'Predictions': y_pred}

    if r2 > best_r2_score:
        best_r2_score = r2
        best_model_name = name

# Save the best model
best_model = best_models[best_model_name]
with open('best_model.pkl', 'wb') as file:
    pickle.dump(best_model, file)

# Print the results
for name, metrics in results.items():
    print(f"Model: {name}")
    print(f"R-squared: {metrics['R_2 Score']}")
    print(f"Mean Squared Error: {metrics['Mean Squared Error']}\n")

print(f"Best model: {best_model_name} with R-squared:
{best_r2_score}")

```

Tuning Linear Regression...

```

/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/
_search.py:317: UserWarning: The total space of parameters 8 is
smaller than n_iter=10. Running 8 iterations. For exhaustive searches,
use GridSearchCV.
    warnings.warn(

```

Best parameters for Linear Regression: {'positive': True, 'fit_intercept': False, 'copy_X': True}

Best R-squared for Linear Regression: 0.3612436852658483

Tuning Random Forest Regression...

```

/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/
_validation.py:528: FitFailedWarning:
25 fits failed out of a total of 50.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.

```

Below are more details about the failures:

```

-----
25 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/_validation.py", line 866, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.12/dist-packages/sklearn/base.py", line 1382, in wrapper
    estimator._validate_params()
  File "/usr/local/lib/python3.12/dist-packages/sklearn/base.py", line 436, in _validate_params
    validate_parameter_constraints(
  File
"/usr/local/lib/python3.12/dist-packages/sklearn/utils/_param_validation.py", line 98, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of RandomForestRegressor must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'sqrt', 'log2'} or None. Got 'auto' instead.

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/_search.py:1108: UserWarning: One or more of the test scores are non-finite:
[0.74303286 0.71414601 0.7550262          nan          nan          nan
          nan          nan 0.73735684 0.7561926 ]
    warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/base.py:1389:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples,), for
example using ravel().
    return fit_method(estimator, *args, **kwargs)

Best parameters for Random Forest Regression: {'n_estimators': 200,
'min_samples_split': 5, 'min_samples_leaf': 1, 'max_features': 'sqrt',
'max_depth': 20}
Best R-squared for Random Forest Regression: 0.7561925973532084
Tuning Gradient Boosting Regressor...

/usr/local/lib/python3.12/dist-packages/sklearn/ensemble/_gb.py:672:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples,), for
example using ravel().
    y = column_or_1d(y, warn=True) # TODO: Is this still required?

Best parameters for Gradient Boosting Regressor: {'subsample': 0.9,
'n_estimators': 500, 'min_samples_split': 5, 'min_samples_leaf': 4,
'max_depth': 5, 'learning_rate': 0.05}

```

Best R-squared for Gradient Boosting Regressor: 0.7729398486349339
Tuning XGBoost Regressor...

Best parameters for XGBoost Regressor: {'subsample': 0.9,
'n_estimators': 200, 'max_depth': 5, 'learning_rate': 0.05,
'colsample_bytree': 1.0}

Best R-squared for XGBoost Regressor: 0.7714925289154053

Model: Linear Regression

R-squared: 0.575990210280537

Mean Squared Error: 471459639.6233815

Model: Random Forest Regression

R-squared: 0.8509345793150533

Mean Squared Error: 165746950.23652834

Model: Gradient Boosting Regressor

R-squared: 0.8533296309538768

Mean Squared Error: 163083874.50125778

Model: XGBoost Regressor

R-squared: 0.7927759885787964

Mean Squared Error: 230413936.0

Best model: Gradient Boosting Regressor with R-squared:
0.8533296309538768

!pip install pandas xgboost scikit-learn pyqt5

Requirement already satisfied: pandas in
/usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: xgboost in
/usr/local/lib/python3.12/dist-packages (3.0.5)
Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: pyqt5 in
/usr/local/lib/python3.12/dist-packages (5.15.11)
Requirement already satisfied: numpy>=1.26.0 in
/usr/local/lib/python3.12/dist-packages (from pandas) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: nvidia-nccl-cu12 in
/usr/local/lib/python3.12/dist-packages (from xgboost) (2.27.3)
Requirement already satisfied: scipy in

```
/usr/local/lib/python3.12/dist-packages (from xgboost) (1.16.2)
Requirement already satisfied: joblib>=1.2.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: PyQt5-sip<13,>=12.15 in
/usr/local/lib/python3.12/dist-packages (from pyqt5) (12.17.1)
Requirement already satisfied: PyQt5-Qt5<5.16.0,>=5.15.2 in
/usr/local/lib/python3.12/dist-packages (from pyqt5) (5.15.17)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2-
>pandas) (1.17.0)
```

###Real-time Predictions:

```
import sys
import pandas as pd
import pickle
from PyQt5.QtWidgets import QApplication, QWidget, QLabel, QLineEdit,
QComboBox, QPushButton, QVBoxLayout, QHBoxLayout, QMessageBox

# Load the trained model
with open('best_model.pkl', 'rb') as file:
    best_model = pickle.load(file)

# Function to convert user input features to numerical values based on
label encoding mappings
def convert_features_to_numerical(features):
    # Label encoding mappings
    company_mapping = {'Acer': 0, 'Apple': 1, 'Asus': 2, 'Chuwi': 3,
'Dell': 4, 'Fujitsu': 5,
                        'Google': 6, 'HP': 7, 'Huawei': 8, 'LG': 9,
'Lenovo': 10, 'MSI': 11,
                        'Mediacom': 12, 'Microsoft': 13, 'Razer': 14,
'Samsung': 15, 'Toshiba': 16,
                        'Vero': 17, 'Xiaomi': 18}
    typename_mapping = {'2 in 1 Convertible': 0, 'Gaming': 1,
'Netbook': 2, 'Notebook': 3,
                        'Ultrabook': 4, 'Workstation': 5}
    ops_mapping = {'No OS': 0, 'Windows': 1, 'macOS': 2,
'other/Linux': 3}
    touchscreen_mapping = {'NonTouchscreen': 0, 'Touchscreen': 1}
    cpu_series_mapping = {'AMD A10-Series': 0, 'AMD A12-Series': 1,
'AMD A6-Series': 2, 'AMD A9-Series': 3,
                        'AMD E-Series': 4, 'AMD FX': 5, 'AMD Ryzen':
6, 'Intel Atom': 7, 'Intel Celeron': 8,
                        'Intel Core M': 9, 'Intel Core i3': 10,
```

```

'Intel Core i5': 11, 'Intel Core i7': 12,
                                'Intel Pentium': 13, 'Intel Xeon': 14,
'Other': 15}
    gpu_category_mapping = {'AMD': 0, 'ARM Mali': 1, 'Intel': 2,
'Nvidia': 3}
    panel_type_mapping = {'4K Ultra HD': 0, 'Full HD': 1, 'IPS Panel':
2, 'Quad HD+': 3, 'Unknown': 4}

    # Convert "Inches" to numeric
    features['Inches'] = pd.to_numeric(features['Inches'],
errors='coerce')

    # Convert features to numerical values based on mappings
    features['Company'] = company_mapping.get(features['Company'], -1)
    features['TypeName'] = typename_mapping.get(features['TypeName'],
-1)
    features['OS'] = ops_mapping.get(features['OS'], -1)
    features['Touchscreen'] =
touchscreen_mapping.get(features['Touchscreen'], -1)
    features['Cpu_Series'] =
cpu_series_mapping.get(features['Cpu_Series'], -1)
    features['GpuCategory'] =
gpu_category_mapping.get(features['GpuCategory'], -1)
    features['PanelType'] =
panel_type_mapping.get(features['PanelType'], -1)

    return features

# Function to make predictions
def predict_price(model, data):
    # Make predictions
    prediction = model.predict(data)
    return prediction

# PyQt5 application
class PricePredictorApp(QWidget):
    def __init__(self):
        super().__init__()
        self.initUI()

    def initUI(self):
        self.setWindowTitle('Laptop Price Predictor')

        layout = QVBoxLayout()

        self.company_label = QLabel('Company')
        self.company = QComboBox()
        self.company.addItem(['Acer', 'Apple', 'Asus', 'Chuwi',
'Dell', 'Fujitsu', 'Google', 'HP', 'Huawei', 'LG', 'Lenovo', 'MSI',
'Mediacom', 'Microsoft', 'Razer', 'Samsung', 'Toshiba', 'Vero',

```

```

'Xiaomi']]
    layout.addWidget(self.company_label)
    layout.addWidget(self.company)

    self.type_name_label = QLabel('Type Name')
    self.type_name = QComboBox()
    self.type_name.addItem('2 in 1 Convertible', 'Gaming',
'Netbook', 'Notebook', 'Ultrabook', 'Workstation'])
    layout.addWidget(self.type_name_label)
    layout.addWidget(self.type_name)

    self.ram_label = QLabel('RAM (GB)')
    self.ram = QLineEdit()
    layout.addWidget(self.ram_label)
    layout.addWidget(self.ram)

    self.inches_label = QLabel('Screen Size (Inches)')
    self.inches = QLineEdit()
    layout.addWidget(self.inches_label)
    layout.addWidget(self.inches)

    self.os_label = QLabel('Operating System')
    self.os = QComboBox()
    self.os.addItem('other/Linux', 'No OS', 'Windows', 'macOS'])
    layout.addWidget(self.os_label)
    layout.addWidget(self.os)

    self.cpu_series_label = QLabel('CPU Series')
    self.cpu_series = QComboBox()
    self.cpu_series.addItem('AMD A10-Series', 'AMD A12-Series',
'AMD A6-Series', 'AMD A9-Series', 'AMD E-Series', 'AMD FX', 'AMD
Ryzen', 'Intel Atom', 'Intel Celeron', 'Intel Core M', 'Intel Core
i3', 'Intel Core i5', 'Intel Core i7', 'Intel Pentium', 'Intel Xeon',
'Other'])
    layout.addWidget(self.cpu_series_label)
    layout.addWidget(self.cpu_series)

    self.touchscreen_label = QLabel('Touchscreen')
    self.touchscreen = QComboBox()
    self.touchscreen.addItem('NonTouchscreen', 'Touchscreen'])
    layout.addWidget(self.touchscreen_label)
    layout.addWidget(self.touchscreen)

    self.gpu_category_label = QLabel('GPU Category')
    self.gpu_category = QComboBox()
    self.gpu_category.addItem('AMD', 'ARM Mali', 'Intel',
'Nvidia'])
    layout.addWidget(self.gpu_category_label)
    layout.addWidget(self.gpu_category)

```

```

self.ssd_label = QLabel('SSD (GB)')
self.ssd = QLineEdit()
layout.addWidget(self.ssd_label)
layout.addWidget(self.ssd)

self.hdd_label = QLabel('HDD (GB)')
self.hdd = QLineEdit()
layout.addWidget(self.hdd_label)
layout.addWidget(self.hdd)
self.flash_storage_label = QLabel('Flash Storage (GB)')
self.flash_storage = QLineEdit()
layout.addWidget(self.flash_storage_label)
layout.addWidget(self.flash_storage)

self.hybrid_label = QLabel('Hybrid (GB)')
self.hybrid = QLineEdit()
layout.addWidget(self.hybrid_label)
layout.addWidget(self.hybrid)

self.panel_type_label = QLabel('Panel Type')
self.panel_type = QComboBox()
self.panel_type.addItem('4K Ultra HD', 'Full HD', 'IPS
Panel', 'Quad HD+', 'Unknown'])
layout.addWidget(self.panel_type_label)
layout.addWidget(self.panel_type)

self.x_res_label = QLabel('X Resolution')
self.x_res = QLineEdit()
layout.addWidget(self.x_res_label)
layout.addWidget(self.x_res)

self.y_res_label = QLabel('Y Resolution')
self.y_res = QLineEdit()
layout.addWidget(self.y_res_label)
layout.addWidget(self.y_res)

self.predict_button = QPushButton('Predict')
self.predict_button.clicked.connect(self.on_predict)
layout.addWidget(self.predict_button)

self.setLayout(layout)

def on_predict(self):
    try:
        input_data = {
            'Company': self.company.currentText(),
            'TypeName': self.type_name.currentText(),
            'Ram': float(self.ram.text()),
            'Inches': float(self.inches.text()),
            'OS': self.os.currentText(),

```



```

        'Cpu_Series': self.cpu_series.currentText(),
        'Touchscreen': self.touchscreen.currentText(),
        'GpuCategory': self.gpu_category.currentText(),
        'SSD': int(self.ssd.text()),
        'HDD': int(self.hdd.text()),
        'Flash Storage': int(self.flash_storage.text()),
        'Hybrid' : int(self.hybrid.text()),
        'PanelType': self.panel_type.currentText(),
        'X_res': int(self.x_res.text()),
        'Y_res': int(self.y_res.text())
    }

    input_df =
pd.DataFrame([convert_features_to_numerical(input_data)])
    prediction = predict_price(best_model, input_df)
    QMessageBox.information(self, "Predicted Price", f'$
{prediction[0]:,.2f}')
    except Exception as e:
        QMessageBox.critical(self, "Error", str(e))

def main():
    app = QApplication(sys.argv)
    ex = PricePredictorApp()
    ex.setGeometry(100, 100, 800, 600) # Set window geometry (x_pos,
y_pos, width, height)
    ex.show()
    sys.exit(app.exec_())

if __name__ == '__main__':
    main()

```